
GPT-Powered Private Knowledge Base

Description

A self-hosted web-based commercial knowledge base for your proprietary content, built on OpenAI GPT-3/4, with support for ongoing ChatGPT-style conversations.

Features

- Unlimited number of separate knowledge bases.
- Up to 100,000 total reference paragraphs (10-20 million words).
- Administrator panel for managing users, content and settings.
- Manage users separately for each knowledge base.
- Add content by copy-paste, URL or PDF/DOC/DOCX/TXT upload.
- Automatic paragraph extraction and semantic analysis of added content.
- Browse content by date, keyword search and semantic search.
- Edit and delete existing content.
- Customizable web form for interrogating knowledge base. Response includes references used and interface allows ongoing conversations.
- Optional per-user and per-IP daily limits on usage of web form.
- JSON-RPC API endpoint for interrogating knowledge base. Response includes references used and ID to allow ongoing conversations.
- Optional per-key daily limits on usage of JSON-RPC API.
- All requests, responses and conversations are recorded.
- Browse recorded interactions by date, user, keyword search and semantic search.

Server requirements

- On-premise or cloud-based Linux server (supports AWS, Azure, GCP and more).
- Ubuntu 12.04+, CentOS 6.2+, Debian 7+, Fedora 15+ or RHEL 6.2+.
- 4 GB RAM + 40 GB SSD drive.

Software requirements

- Apache 2
- PHP 5+
- MySQL/MariaDB 5.5+

Licensing and pricing

- Free 30-day demonstration.
- Fixed annual license fee (includes installation support and updates).
- Customer pays OpenAI directly for GPT-3/4 API usage.

Scheduled delivery

- May 2023: Alpha release.
- June 2023: Beta release.
- July 2023: General release.